



THE GROUNDWORK FIELD GUIDE · VERSION 1.1 · BY ADAM ABDALLA

THE GROUNDWORK FIELD GUIDE

*The mandatory foundation before agents. Identity, Knowledge,
Governance, Continuity, Measurement, and Operating Rhythm.
Then one Operator.*

[GROUNDWORK.ADAMABDALLA.COM](https://groundwork.adamabdalla.com)

DO NOT START WITH AGENTS

SIX LAYERS · FIFTEEN DOCUMENTS

HOW TO USE THIS BOOK

This is the whole program on paper: the same doctrine as groundwork.adamabdalla.com, in printable form. Build the documents in order. Activate the Operator only after Phase 1 exists. Do not start with a swarm.

- | | |
|-----------|-------------------------------------|
| 01 | What this is |
| 02 | The six layers |
| 03 | The build path |
| 04 | Playbook |
| 05 | Readiness scorecard |
| 06 | The fifteen documents |
| 07 | The Operator context package |
| 08 | The Operator prompt |
| 09 | Agent scope and multi-agent |

WHAT THIS IS

Groundwork is not a promise. It is a folder of files. You build fifteen plain-text documents in a fixed order, and those documents become the foundation any AI can stand on.

Most teams skip this and go straight to agents. The result is generic output, invented policy, and systems that cannot tell doctrine from a draft. Groundwork installs the missing layer: who you are, what your system may trust, what it may do, what survives failure, how you know it is working, and how it is maintained.

The order is not optional, because each layer depends on the one below it. If your identity is unclear, your knowledge is unclassified, and your rules are unwritten, an agent will amplify the mess faster than you can catch it. The rule underneath everything: **do not automate chaos**.

Phase 1 first

Identity, Knowledge, and Governance.
Eleven documents, required first. Output quality and decision boundaries improve immediately, before a single agent exists.

Then one Operator

A single scoped agent. Not autonomous. Not a swarm. It drafts, cites authority, and stops when it should.

Then durability

Continuity, Measurement, and Operating Rhythm: four documents that make the foundation survive time, people, and drift.

Multi-agent last

Optional. Only after the Operator is stable. Duplicate the folder pattern with a narrower charter. The human still decides who is invoked.

Everything here is free and ungated: this book, the website, and the open curriculum on GitHub. The files are yours, in plain text, on your machine. Markdown works with Claude, ChatGPT, Gemini, or any future AI. No lock-in. Ever.

THE SIX LAYERS

Six layers, built in order, across three phases. Phase 1 is Identity, Knowledge, and Governance. The rest follows once the Operator is live.

LAYER	PHASE	WHAT IT SETTLES
01 Identity	Phase 1	Who you are and how you are allowed to behave.
02 Knowledge	Phase 1	What the system is allowed to trust and use.
03 Governance	Phase 1	What the system is allowed to do.
04 Continuity	Phase 3	What survives over time, people, and failure.
05 Measurement	Phase 3	Whether the foundation is actually working.
06 Operating Rhythm	Phase 3	How the foundation is maintained so it does not decay.

Phase 2 sits between them and is not a layer: it is the activation of one scoped Operator on top of the eleven Phase 1 documents. Chapter 07 covers it.

The test for every layer: could a stranger read its documents and then write, decide, and refuse the way you would? If yes, the layer is done. If no, keep sharpening.

THE BUILD PATH

You do not need to understand the whole system to start. You need one sitting, one folder, and one honest page.

FIRST 30 MINUTES

Day one

Take the readiness scorecard and write down your weakest layer. Read the four building blocks in Foundations on the site: the model is a well-read librarian, agents are its hands, prompts are instructions, and your documents are the fuel. Then create a folder called `Groundwork`, inside it `01_Identity`, and write a rough Identity Charter. A rough draft counts. That is day one, done: you know your gap, you understand the machine, and the first of fifteen documents exists on your machine, owned by you.

PHASE 1 · A FOCUSED WEEKEND

Foundation install

Eleven documents. Output quality and decision boundaries improve immediately.

- Identity Charter, Brand Voice, Decision Principles
- Authority Classification, Context Vault, Strategic Priorities, Active Constraints
- Automation Decision Framework, Escalation Rules, Allowed Tools, Agent Scope

PHASE 2 · ABOUT A WEEK

First Operator live

One scoped agent, running real work inside the foundation.

- Lock the Agent Scope
- Assemble the Operator context package
- Run real Draft & Wait work through it
- Correct the documents from actual misses

PHASE 3 · AS REQUIRED

Durability

The foundation survives time, people, and drift.

- Continuity & Succession
- Session Notes discipline
- Measurement Notes and the readiness scorecard
- Operating Rhythm

PLAYBOOK

Your first working session

1. Name the workspace after your business. One folder, plain text inside.
2. Score the six layers out loud, honestly. Write notes on the gaps.
3. Only the required documents come first. Resist building ahead.
4. Write the Identity Charter in one honest sitting.
5. If time remains, lock Brand Voice and Authority Classification.
6. Do not promise yourself agents. Build a foundation the Operator can use.

When to activate the Operator

After all eleven Phase 1 documents exist and a named human has signed off on the governance layer. Confirm the Agent Scope. Load the package. Run one real Draft & Wait task. Then correct the documents.

Multi-agent

Do not build a swarm on a weak foundation. If a second role is later needed, duplicate the folder pattern with a narrower charter. The human decides who is invoked.

Do not

- Invent policy
- Treat drafts as doctrine
- Add unauthorized tools
- Promise autonomy
- Skip escalation rules

READINESS SCORECARD

Score each layer 0 to 5, out loud, honestly. Notes go in the last column.

Phase 1 must be honest before an Operator is activated.

LAYER	QUESTION	SCORE 0-5	NOTES
01 Identity	Would a stranger write like us after reading this?		
02 Knowledge	Can the system tell doctrine from a draft?		
03 Governance	Does it know when to act, draft, or stop?		
04 Continuity	Would this survive the founder being gone for 90 days?		
05 Measurement	How would we know this is working next month?		
06 Rhythm	Who owns each layer, and when is it reviewed?		

0 = absent. 3 = usable. 5 = another operator could inherit it without a briefing.

This is the one-page meeting version. The full instrument, 24 weighted questions scored in your browser with nothing collected or transmitted, runs free at groundwork.adamabdalla.com/guide/scorecard.html.

THE FIFTEEN DOCUMENTS

Every document in the system, with its file path, its template, and the standard for writing it well. Copied straight from the guide pages, which remain the living version.

LAYER ONE · PHASE 1 · REQUIRED FIRST

Identity.

Who you are and how you are allowed to behave. Three documents that make every AI answer sound like you, and decide like you.

DOCUMENT 01

Identity Charter

01_Identity/Identity_Charter.md

A non-negotiable statement of who you are, and will never become. One clear paragraph, then two short lists. No marketing language. Describe the actual operator or organization, the way you would to a sharp new hire on day one.

```
# Identity Charter

## Who we are
[One clear paragraph. No marketing language.
Describe the actual operator or organization.]

## What we stand for
-
-
-

## What we will never become
-
-
-

This charter is non-negotiable. All output and
recommendations must remain consistent with it.
```

How to write it well: the “never become” list is the sharp edge. “We will never pressure a tenant into terms we would not accept ourselves” does more governing work than any mission statement. Write at least three, and make each one something you have actually refused.

DOCUMENT 02

Brand Voice

01_Identity/Brand_Voice.md

Exact rules for how the system writes and speaks. Not adjectives: rules. “Professional but friendly” is useless to a machine. “Lead with the point, no hype language, no unnecessary softening” is enforceable.

```
# Brand Voice

## Tone
- Direct
- Calm
- Precise
- Confident without arrogance

## Style rules
- Lead with the point
- Prefer short to medium sentences
- No hype language
- No unnecessary softening

## Vocabulary
Prefer: [words and phrases you actually use]
Avoid:  [words that never sound like you]

## Examples
On-voice: " "
Off-voice: " "
```

How to write it well: the examples carry most of the weight. Paste in two real emails you were proud of and one AI draft that made you cringe, and label them. A model learns more from one on-voice/off-voice pair than from ten adjectives.

DOCUMENT 03

Decision Principles

01_Identity/Decision_Principles.md

The five to nine rules that govern judgment when the answer is not obvious. These are the tiebreakers the system reaches for when two good options conflict.

```
# Decision Principles
```

1. Clarity over cleverness
2. Truth over comfort
3. Long-term over short-term
4. Authority must be explicit
5. Reversibility matters
6. Context before action

Add or rewrite these so they match how
you actually decide.

How to write it well: do not copy the defaults. Recall three hard calls you made in the last year and ask what rule you were following. If a principle would not have changed a real decision, it is a poster, not a principle.

Knowledge.

What the system is allowed to trust and use. Four documents that stop your AI from treating a rumor and a contract as equal.

DOCUMENT 04

Authority Classification

03_Authority_Rules/Authority_Classification.md

```
# Authority Classification

## Level 0: Core Doctrine
Non-negotiable. Overrides everything.
Examples: Identity Charter, Decision Principles,
legal mandates, operating agreements.

## Level 1: Institutional Standards
Approved ways of working.
Examples: Brand Voice, internal policies,
approval thresholds, communication standards.

## Level 2: Validated Knowledge
Reviewed and accepted.
Examples: executed contracts, cleaned data,
completed reports, verified third-party info.

## Level 3: Working Material
Useful but not authoritative.
Examples: drafts, preliminary analysis,
internal notes, external proposals.

## Level 4: Untrusted / Raw
Requires caution.
Examples: forwarded emails, unverified quotes,
market chatter, personal notes.

Rule: Level 0 outranks everything: the lower the number, the higher the authority.
Cite the Authority Level when making material
recommendations.
```

How to apply it: you do not need to label every file you own on day one. Classify as you touch things. When a document enters the vault or gets used in a real decision, give it a level: a line of frontmatter or a single word in the filename is enough.

DOCUMENT 05

Context Vault v1

02_Current_Context/Context_Vault.md

The short list of files the system must always load first. This is not your archive. It is the briefing packet.

```
# Context Vault v1

## Always load
- Identity Charter
- Brand Voice
- Decision Principles
- Current Strategic Priorities
- Active Constraints

## Domain files (add only what is necessary)
-
-
-

Keep this list short. A vault that contains
everything contains nothing.
```

How to apply it: the discipline is subtraction. Every file on the always-load list costs attention on every single task. If a file has not changed an answer in a month, demote it to a domain file that gets loaded on request.

DOCUMENT 06

Strategic Priorities

02_Current_Context/Strategic_Priorities.md

What is actually being optimized this cycle, and what is not. Without this file, the system optimizes for whatever the last prompt implied.

```
# Strategic Priorities
Last updated: [date]

## Current top priorities (in order)
1.
2.
3.

## Active focus this cycle
-

## Explicit non-priorities
- [things you are deliberately not doing,
  so the system stops suggesting them]
```

How to apply it: the non-priorities section works harder than the priorities. “No new acquisitions this quarter” kills a whole category of well-meaning but useless AI suggestions. Date the file and re-touch it whenever the cycle turns.

DOCUMENT 07

Active Constraints

02_Current_Context/Active_Constraints.md

Hard limits, soft limits, and current capacity. Constraints are context: an AI that does not know your limits will recommend plans you cannot run.

```
# Active Constraints

## Hard constraints
- No irreversible external actions without
  human approval
- No access to financial systems
- No legal document finalization
- No public communications without review

## Soft constraints
- Prefer internal tools over new third-party
  services
- Keep recommendations tightly scoped
- Prefer short drafts that need minimal revision

## Current capacity notes
-
```

How to apply it: hard constraints are enforced in Layer 3 governance; here they are stated as fact. Capacity notes are the human part: “one person, ten hours a week for this” changes every recommendation the system should make.

Governance.

What the system is allowed to do. Four documents that decide, in writing and before the fact, when AI acts, when it drafts, and when it must stop.

DOCUMENT 08

Automation Decision Framework

03_Authority_Rules/Automation_Framework.md

```
# Automation Decision Framework

## Zone 1: Auto-Execute
Criteria: reversible, internal, no exposure.
Tasks:
-
-

## Zone 2: Draft & Wait
Criteria: external-facing or judgment-adjacent.
Approver: [named human]
Tasks:
-
-

## Zone 3: Never Automate
Criteria: irreversible, legal, financial,
precedent-setting.
Tasks:
-
-

Default: any task not listed is Zone 2 until
a human assigns it a zone.
```

How to write it well: start by listing the ten tasks you actually want help with, then zone each one. The default rule at the bottom matters most: unlisted work falls to Draft & Wait, never to auto-execute.

DOCUMENT 09

Escalation Rules

03_Authority_Rules/Escalation_Rules.md

When the system must stop and hand the decision to a human, even in the middle of an approved task.

```
# Escalation Rules

Stop and escalate to [named human] when:
- Confidence is low or sources conflict
- Authority levels conflict (e.g., a draft
  contradicts doctrine)
- The task drifts outside its assigned zone
- Money, legal exposure, or a relationship
  is at stake beyond stated thresholds
- Anyone (including the requester) asks to
  bypass these rules

When escalating: state what was being done,
what triggered the stop, and what decision
is needed. Do not proceed on silence.
```

How to write it well: the last line is the one that matters. “Do not proceed on silence” converts every ambiguous situation into a paused one instead of a guessed one.

DOCUMENT 10

Allowed Tools & Access

04_Capabilities/Allowed_Tools.md

The explicit inventory of what the system may touch, and the systems it may never enter.

```
# Allowed Tools & Access

## May use
- [AI assistant / platform]
- [document drive: read]
- [calendar: read]

## May use with approval per instance
- [email: send]
- [external posting]

## May never access
- Banking and payment systems
- Signature authority
- [anything else that is yours alone]
```

How to write it well: write the “never” list first. It is shorter, and it is the one that protects you. Everything not on a list is off-limits by default.

DOCUMENT 11

Agent Scope (First Operator)

04_Capabilities/Agent_Scope.md

The job description for your first Operator, written before the Operator exists. One role, a short task list, and hard boundaries.

```
# Agent Scope: First Operator

## Role
[One sentence. e.g., "Drafting and document
operator for tenant communications."]

## In scope (Zone 1 and Zone 2 tasks only)
-
-
-

## Out of scope
- Everything in Zone 3
- Any task not listed above

## Operating rules
- Cite Authority Levels for material claims
- Follow Escalation Rules without exception
- All Zone 2 output goes to [named human]
```

How to write it well: narrow wins. “Handles tenant communication drafts” produces a competent operator. “Helps with the business” produces an improviser. You can widen scope later, deliberately, one addition at a time.

Continuity.

What survives over time, people, and failure. Two documents that make the foundation outlive the founder being busy, absent, or gone.

DOCUMENT 12

Continuity & Succession

01_Identity/Continuity.md

How the foundation survives people leaving and systems failing. It lives in the Identity folder on purpose: what survives you is part of who you are.

```
# Continuity & Succession

## Key context that must survive
- Identity Charter
- Decision Principles
- Authority Classification
- Current Strategic Priorities
- Critical operating files

## Succession rules
- Who inherits decision rights:
- How access is transferred:
- What must be reviewed before transfer:

## Failure & override
- How the system is stopped:
- Who has final override authority:
- What gets preserved during failure:
```

How to write it well: use names, not roles-you-hope-to-hire. “Your brother gets read access and your attorney gets the override call” is a real succession rule. Then test the boring parts: does the successor actually have the passwords, and do they know the vault exists?

DOCUMENT 13

Session Notes

05_Working_Memory/Session_Notes.md

Temporary working memory, kept deliberately separate from doctrine. This is the file that stops two kinds of rot: losing decisions made mid-session, and letting session chatter quietly harden into policy.

```
# Session Notes
*Temporary working memory only. Clear or archive
at the end of major sessions.*

## Current session
- Date:
- Primary objective:
- Key decisions made:
- Open questions:
- Items requiring human input:
```

How to write it well: the discipline is the clearing. At the end of a significant session, anything worth keeping gets promoted into a real document at a real authority level, and the notes get archived or wiped. Working memory that never gets cleared becomes a second, contradictory source of truth.

Measurement.

Whether the foundation is actually working. One document and a monthly habit that catch drift before drift becomes damage.

DOCUMENT 14

Measurement Notes

06_Output_Standards/Measurement_Notes.md

How you will know the foundation is working, written down so the check actually happens.

```
# Measurement Notes

## Signals of a working foundation
- Output stays on-voice without prompting
- The system cites authority correctly
- Escalations are appropriate, not constant
- Humans trust the drafts enough to use them

## Review questions (monthly)
- What broke this month?
- What did the system get wrong?
- What context was missing?
- What should be promoted or demoted
  in authority?
```

How to use it: the four signals are observations, not metrics dashboards. You can answer all four from a week of normal use plus your approval log. Write one dated entry per month under the questions. Six entries in, you have something rare: an evidence trail of whether your AI system is getting better or worse.

Operating Rhythm.

How the foundation is maintained so it does not decay. One document that turns fourteen files into a system that stays true.

DOCUMENT 15

Operating Rhythm

06_Output_Standards/Operating_Rhythm.md

Ownership, cadence, and how changes get made. The last document in the program, and the one that keeps the other fourteen alive.

```
# Operating Rhythm

## Ownership
- Identity:
- Knowledge:
- Governance:
- Continuity:
- Measurement:

## Review cadence
- Weekly:
- Monthly:
- Quarterly:

## Update protocol
- How changes are proposed:
- Who approves:
- How versioning is handled:
```

How to write it well: in a small operation, every line may say the same name, and that is fine. The point is that it is written, so the day someone else joins, ownership transfers as an edit instead of an excavation.

THE OPERATOR CONTEXT PACKAGE

One scoped agent, activated only after Phase 1 exists. Not autonomous. Not a swarm. A competent internal operator that drafts, cites authority, and stops when it should.

When you are allowed to activate: all eleven Phase 1 documents exist, a named human has signed off on the governance layer, and the Agent Scope file describes one role in one sentence. If any of that is missing, the Operator will improvise it, and improvised policy is how automation goes wrong.

What an Operator is not: it is not autonomous, it does not act while you sleep, it does not have memory you have not given it, and it is not a fleet. It is one agent, one context package, one human in the loop. For most owners this is enough. Permanently.

Everything the Operator knows lives in one folder, loaded in order. This is the package: the fifteen Groundwork documents arranged for loading. A vault that contains everything contains nothing; this is the deliberate subset.

```
01_Operator/  
  00_CONTEXT.md  
  01_Identity/  
    Identity_Charter.md  
    Brand_Voice.md  
    Decision_Principles.md  
    Continuity.md  
  02_Current_Context/  
    Strategic_Priorities.md  
    Active_Constraints.md  
    Context_Vault.md  
  03_Authority_Rules/  
    Authority_Classification.md  
    Automation_Framework.md  
    Escalation_Rules.md  
  04_Capabilities/  
    Allowed_Tools.md  
    Agent_Scope.md  
  05_Working_Memory/  
    Session_Notes.md  
  06_Output_Standards/  
    Measurement_Notes.md  
    Operating_Rhythm.md
```

00_CONTEXT.md is a one-page index: what this package is, the load order, and the one-line rule: higher authority levels override lower ones. Load the files first, in order. Then paste the Operator prompt as the session instruction.

THE OPERATOR PROMPT

This is the session instruction. Do not expand it into a personality. Keep it as rules.

You are now the Operator for this organization.

Load the files in this package in order. Do not invent policy. Do not expand your authority.

Primary directive: raise the quality of thinking and output while protecting standards.

Rules

- Stay on-voice at all times.
- Cite Authority Level on material claims.
- If the request is Auto-Execute, do it.
- If the request is Draft & Wait, produce the complete draft and stop.
- If the request is Never Automate, refuse and escalate.
- If authority is unclear, confidence is low, or the action is irreversible, escalate.

When escalating, state:

1. What was requested
2. Why it cannot proceed
3. What decision or information is needed

Your output should read like it came from a highly competent operator, not an assistant trying to be helpful.

The first two weeks

Activation is a shakedown cruise, not a launch party. Run it like this:

- **Days 1–3: Zone 1 only.** Internal summaries, formatting, file work. You are testing whether the Operator loads context correctly and cites authority levels, not whether it can do your job.
- **Days 4–14: real Draft & Wait work.** Give it the actual emails, notices, and reports on your desk. Approve, reject, and (this is the important part) when a draft misses, fix the *document* that caused the miss, not just the draft.

- **Log every correction in Session Notes.** Three misses with the same root cause means a foundation file needs editing. That loop (miss, trace, edit, retest) is the whole method.

AGENT SCOPE AND MULTI-AGENT

The Operator may

- + Draft work inside Draft & Wait
- + Summarize Level 0–2 material
- + Flag missing or conflicting authority
- + Prepare recommended next actions

The Operator may not

- Invent policy
- Expand its own authority
- Execute irreversible actions
- Spawn additional agents without a human decision

Adding roles later

The default is one. One Operator, one context package, one human in the loop. Add a second role only after the Operator is stable and a second, clearly different job exists: a Researcher that never drafts client-facing work, or a Reviewer that only checks authority and voice.

How to add one: do not build an orchestration engine. Duplicate the folder pattern. Give the new role a narrower charter, its own Allowed Tools, and the same Escalation Rules. The human decides who is invoked.

What to refuse: agents that create other agents. Shared memory without classification. Any role that can execute Never Automate work.

The Operator is working when

- + Drafts arrive on-voice without re-prompting
- + Material claims come with authority levels attached
- + It refuses Zone 3 requests and escalates cleanly, stating what it needs
- + Corrections per week are falling because the documents keep improving
- + You trust it with a category of work you used to do at 11pm



Groundwork. The AI readiness system by Adam Abdalla. Lafayette, Louisiana.

This Field Guide is free and ungated, like everything in the program. The living version of this doctrine is the website; this book is generated from it, so if the two ever disagree, the site is current.

`groundwork.adamabdalla.com · adam@adamabdalla.com`

Version 1.1 · August 2026